

Supervised Learning - I

M. Vazirgiannis



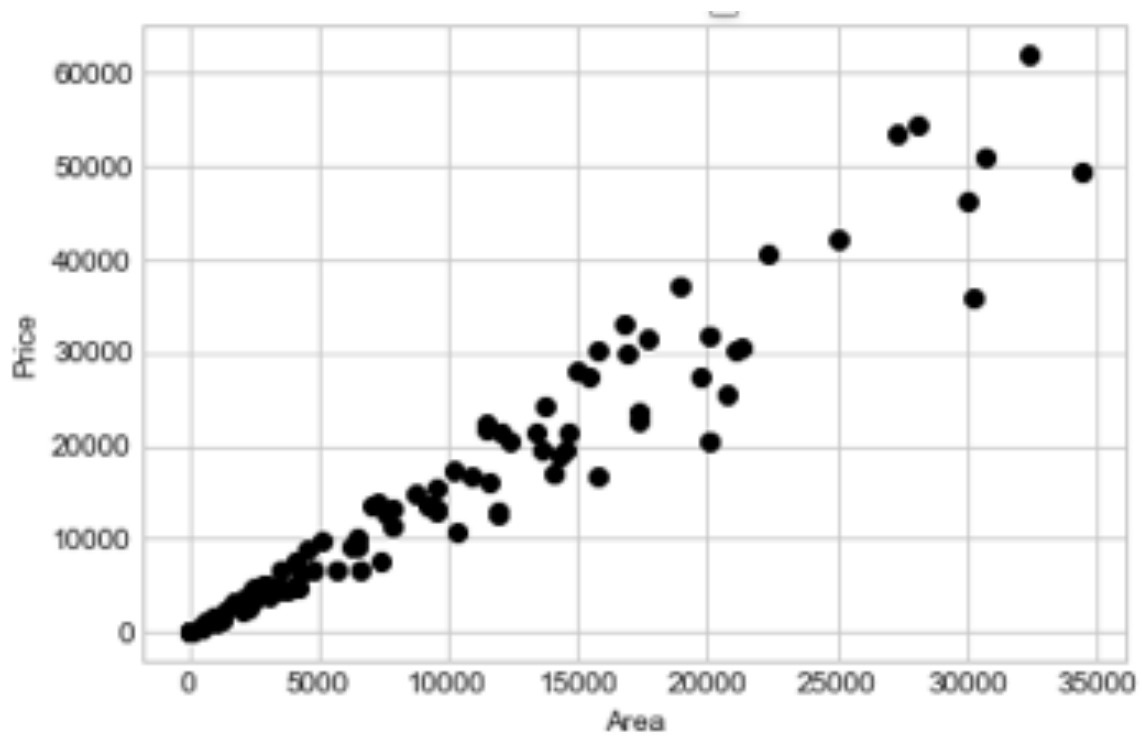
, LIX, École Polytechnique

September 2024

Outline

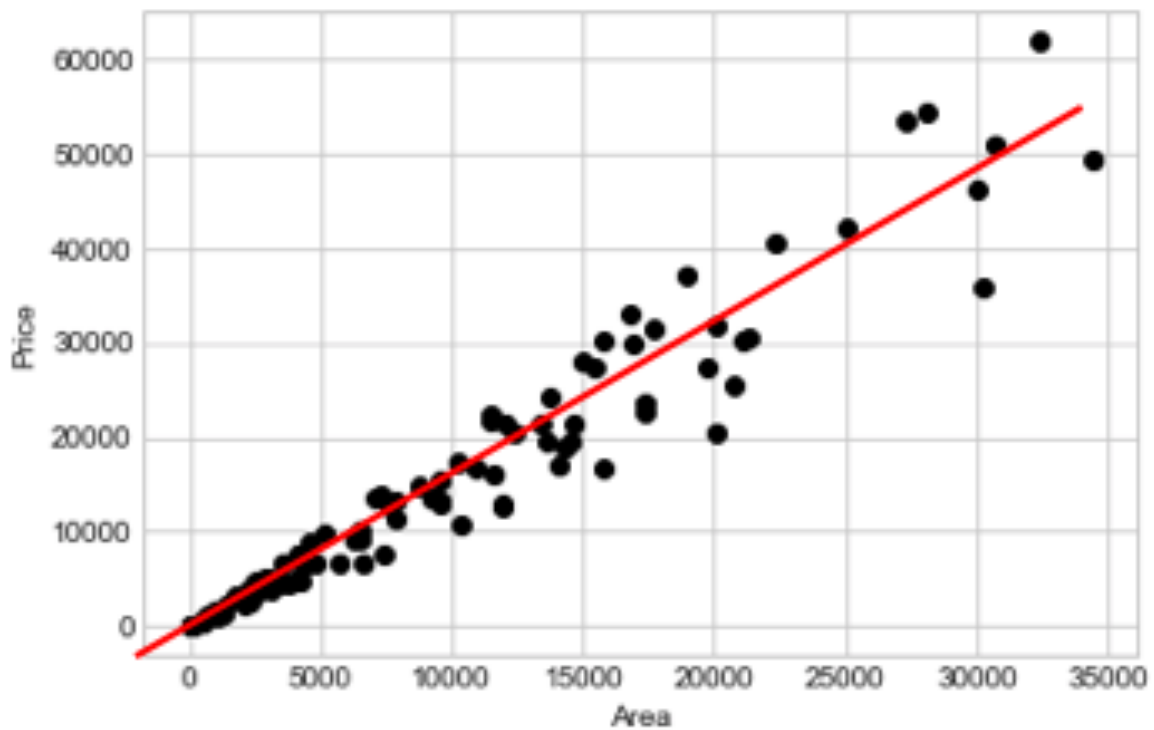
- Introduction to supervised learning
- K-nn
- Naïve Bayes
- Logistic Regression

Supervised Learning



Can we predict the price of a parcel based on its size (surface in m^2) ?

Supervised Learning



Can we predict the price of a parcel based on its size (surface in m^2) ?

Supervised learning

- $X = X_1, X_2, \dots, X_p$: input variables (features)
- Y : “output” or target variable that we try to predict
- A pair (x_i, y_i) is called a training example, where:

$$x_i \in X = X_1 \times \dots \times X_p \text{ and } y_i \in Y$$

- The training set is a list of training examples:

$$\{(x_i, y_i) ; i = 1 \dots m\}$$

- The supervised learning problem is formulated as:
 - Given a training set,
 - Learn a function $h : X \rightarrow Y$ such that $h(x)$ is “good” predictor for the corresponding value of y .

For historical reasons, the function h is called a *hypothesis*.

Supervised Learning – Performance evaluation

Goal: find a hypothesis $h : X \rightarrow Y$ minimizing

$$Error(h) = E_{X,Y} L(Y, h(X))$$

Classification (Y label values, i.e. Y=yes,no) - Error: 0-1 loss:

$$L(Y, h(X)) = 1(Y \neq h(X))$$

Regression Y: number, i.e. Y = stock_price - Error: square loss:

$$L(Y, h(X)) = (Y - h(X))^2$$

Generative vs. Discriminative classifiers

- Generative algorithm

- Data generated by a distribution of feature
- Assume x features of an animal and y the animal genre (i.e cat, dog).
- For unknown classification data x find the class maximizing the posterior $p(y|x)$ based on Bayes rules

$$p(y|x) = \frac{p(x|y)p(y)}{p(x)}$$

- Therefore

$$p(y|x) = \operatorname{argmax}_y(p(x|y)p(y))$$

Generative vs. Discriminative classifiers

- Discriminative algorithm

- Learn directly $p(y|x)$ or learn mappings directly from the space of inputs X to the labels 0, 1
- Do not care about how the data was generated - simply categorizes a given vector
- Generally discriminative classifiers are more effective - better accuracy.

Some Classes of classifiers

- Class-conditional probabilistic based on $p(x | c_k)$
 - **Naïve Bayes**: simple, often effective in high dimensions
 - **Parametric generative models**, e.g., Gaussian (can be effective in low-dimensional problems: leads to quadratic boundaries in general)
- Regression-based $p(c_k | x)$ directly
 - **Logistic regression**: simple
 - **Neural networks**: non-linear extension of logistic regression

Some Classes of classifiers

- Discriminative models, focus on locating optimal decision boundaries
 - **Linear discriminants**: perceptron - simple, sometimes effective.
 - **Support vector machines**: Generalization of linear discriminants, can be quite effective, computational complexity is an issue
 - **Nearest neighbor**: Simple, can scale poorly in high dimensions
 - **Decision trees**: Often effective in high dimensions

Classification: Results & Evaluation

- Confusion matrix

Predicted class	Actual class	
	1	0
1	True Positive	False positive
0	False negative	True negative

- Precision $\frac{TP}{TP+FP}$
- Recall $\frac{TP}{TP+FN}$
- Accuracy $\frac{TP+TN}{TP+TN+FP+FN}$
- $F1 = \frac{2 * \text{precision} * \text{recall}}{\text{precision} + \text{recall}}$

Classification: Results & Evaluation

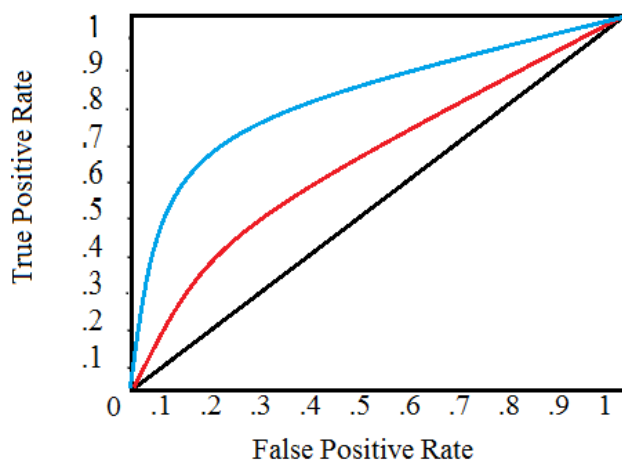
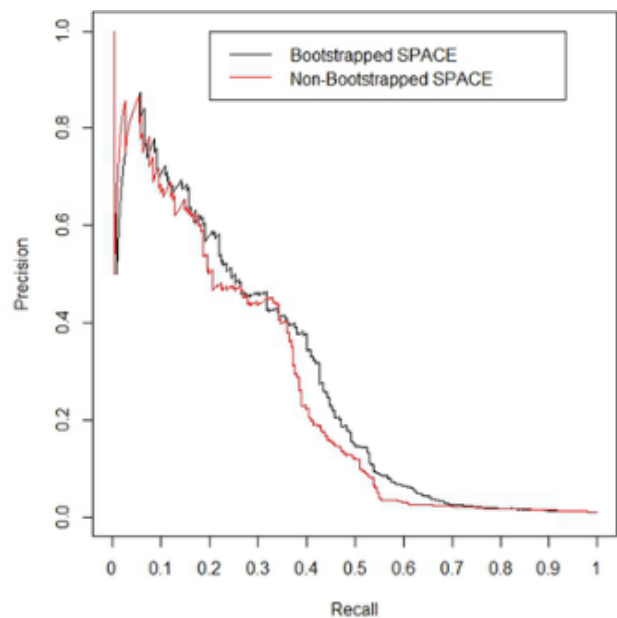


Figure: ROC curve



Precision recall curve

ROC/AUC curve

- True positive rate:

$$TPR = \frac{TP}{TP + FN}$$

- False negative rate:

$$FNR = \frac{FN}{FN + TN}$$

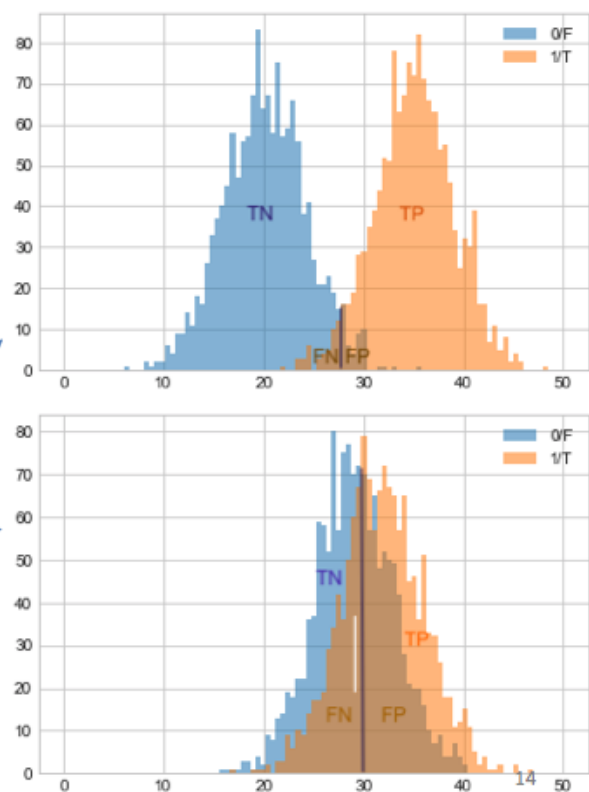
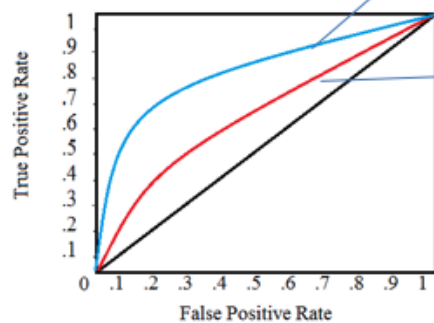


Figure: overlapping of positive / negative classes

Classification: Results & Evaluation

- Log-loss: more refined evaluation of the classification
- Capitalizing on the classifier probability

$$\text{log-loss} = -\frac{1}{N} \sum_{i=1}^N y_i \log p_i + (1 - y_i) \log p_i(1 - p_i)$$

- Log loss = cross entropy $H(p, q) = -\sum_x p(x) \log q(x)$ between the distribution of true labels and predictions
- Closely related to Kullback-Leibler divergence

$$KL(p|q) = \sum_x p(x) \log \frac{p(x)}{q(x)}$$

- Minimizing cross entropy, maximize accuracy of the classifier.

Log Loss Analysis

Advantages of Log Loss

- Interprets Probabilities: directly evaluates how well the predicted probabilities match the true labels, making it ideal for probabilistic models.
- Penalizes Overconfidence: Models overconfident but wrong heavily penalized, encouraging models to be more conservative with their predictions.

Disadvantages of Log Loss

- Sensitive to Imbalanced Classes: for data highly imbalanced, log loss is misleading as it penalizes wrong predictions more than simply misclassifying the majority class.
- Difficult to Interpret: raw value log loss not as intuitive as accuracy or other simple metrics, though very useful for comparing models.

Outline

- Introduction to supervised learning
- **K-nn**
- Naïve Bayes
- Logistic Regression

k -NN Algorithm

- K -NN is an non parametric lazy learning algorithm.
 - No assumptions for the data (i.e. distribution, linearly separable)
 - No training lack of generalization
- K -nn classification: by majority voting
- K -nn regression: average value of k nearest neighbors

K-NN classification example

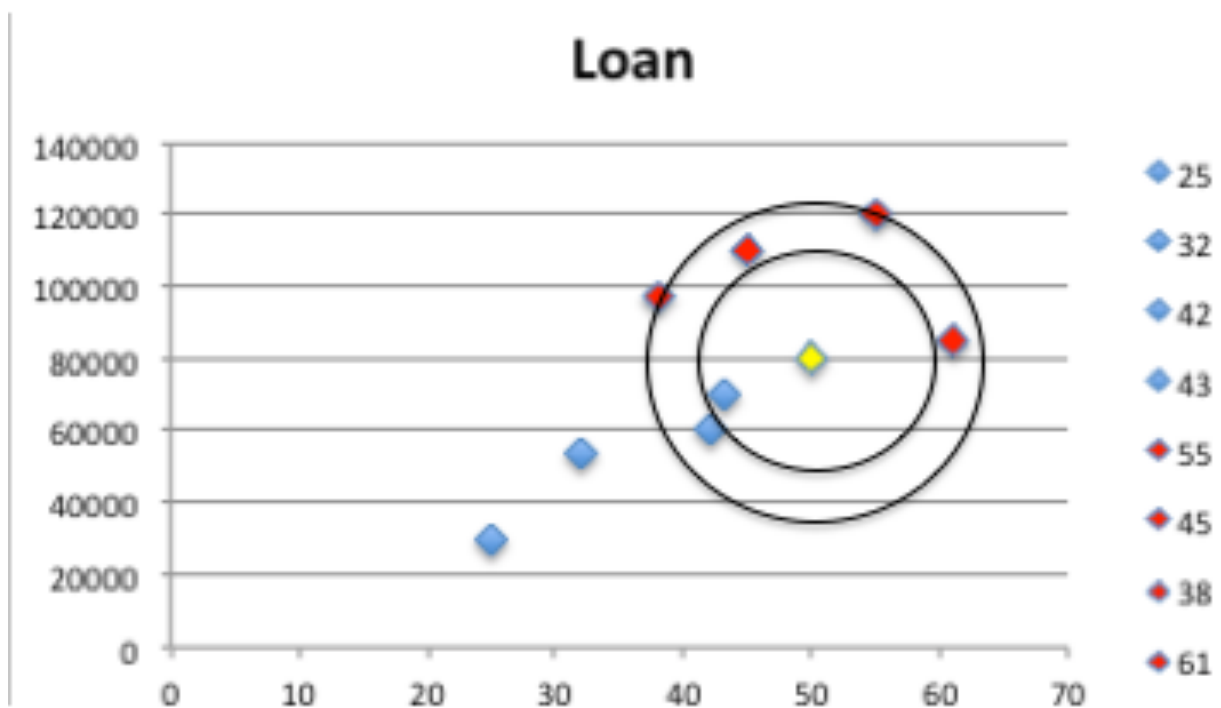


Figure: Loan amount vs. age: decide class by majority of nns

K-NN classification example

Age	Loan	Default	Distance
25	30000	N	50000,0062
32	54000	N	26000,0062
42	60000	N	20000,0016
43	70000	N	10000,0024
55	120000	Y	40000,0003
45	110000	Y	30000,0004
38	97000	Y	17000,0042
61	85000	Y	5000,0121
50	80000	?	

Figure: Loan amount vs. age: decide class by majority of nns

$$D(x, y) = \sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2}$$

K-NN classification example

Age	Loan	Default	Distance
0	0	N	0,88932
0,1944	0,2667	N	0,57746
0,4722	0,3333	N	0,31427
0,5	0,4444	N	0,22395
0,8333	1	Y	0,46564
0,5556	0,8889	Y	0,36111
0,3611	0,7444	Y	0,38313
1	0,6111	Y	0,31056
0,6944	0,5556	?	

Figure: Loan amount vs. age: decide class by majority of nns

$$\frac{X_i - \min(X)}{\max(X) - \min(X)}$$

K-NN regression

- k-NN algorithm to estimate continuous variables.
- weighted average of the k nearest neighbors, weighted by the inverse of their distance:
 - Compute distance of query example to the labeled examples.
 - Order labeled examples by increasing distance.
 - Find a heuristically optimal number k of nearest neighbors, based on RMSE – with cross validation.
 - Calculate inverse distance weighted average with k-nearest multivariate neighbors.

K-NN issues

- K: larger values reduce noise effect but make classes boundaries blur
- Dimensionality affects the performance
- Feature selection/scaling (mutual information)
- Binary classification: k odd number
- popular way for empirically optimal k via bootstrap method
- Assume large number of points and c classes:

$$E_{k-nn} \leq E_{Bayes} \left(2 - \frac{c}{c-1} E_{Bayes} \right)$$

Outline

- Introduction to supervised learning
- K-nn
- **Naïve Bayes**
- Logistic Regression

Bayesian Classification: Why?

- Probabilistic learning
 - Calculate explicit probabilities for hypothesis,
 - practical approaches to certain types of learning problems
- Incremental
 - Each training example can incrementally increase/decrease the probability that a hypothesis is correct.
 - Prior knowledge combined with observed data.
- Probabilistic prediction
 - Predict multiple hypotheses, weighted by their probabilities

Bayesian Classification

- Problem formalized using a-posteriori probabilities:
- $p(C|X)$ = prob. vector $X = \langle x_1, \dots, x_k \rangle$ is class C .
 - e.g. $p(\text{class} = N | \text{outlook} = \text{sunny}, \text{windy} = \text{true}, \dots)$
- Assign to sample X class label C : $p(C|X)$ is maximal
- Bayes theorem:

$$p(C|X) = \frac{p(X|C)p(C)}{p(X)}$$

- $p(X)$: prior probability of vector X
- $p(C)$ = prior probability of class C in training data
- $p(X|C)$ = probability of X given C
- $p(C|X)$ = probability of X given C

Learn probabilistic models with dependence

- Problem: computing $P(X|C)$ not unfeasible! – Why?
- $p(C)$: assume p classes
- $p(X = \langle x_1, \dots, x_k \rangle)$: k binary features
- $p^{(2^k-1)}$ parameters
- Likelihood: $p(X|C)$
 - Need a value for each possible $p(X = \langle x_1, \dots, x_k \rangle | C)$

Naive Bayesian Classification

- Naïve assumption: attribute independence

$$p(X|C) = \prod_i p(x_i|C)$$

- $p(x_i|C)$: relative frequency of x_i as i-th attribute in class C
- Computationally feasible
- Generative probabilistic model with conditional independence assumption
- Prediction:

$$p(C|X) = \operatorname{argmax}_C (p(X|C)p(C)) = \operatorname{argmax}_C \left(\prod_i^k p(x_i|C)p(C) \right)$$

Naive Bayesian Classification

- Simple to train
- estimate conditional probabilities for each feature-class pair
- Often very good baseline
- Feature selection can be helpful, e.g., information gain
- However. . . . on most problems can usually be outperformed by a more complex model

Playtennis example: estimating $P(x_i|C)$

Outlook	Temperature	Humidity	Windy	Class
sunny	hot	high	false	N
sunny	hot	high	true	N
overcast	hot	high	false	P
rain	mild	high	false	P
rain	cool	normal	false	P
rain	cool	normal	true	N
overcast	cool	normal	true	P
sunny	mild	high	false	N
sunny	cool	normal	false	P
rain	mild	normal	false	P
sunny	mild	normal	true	P
overcast	mild	high	true	P
overcast	hot	normal	false	P
rain	mild	high	true	N

$$p(P) = 9/14$$

$$p(N) = 5/14$$

outlook	
$p(\text{sunny} P) = 2/9$	$p(\text{sunny} N) = 3/5$
$p(\text{overcast} P) = 4/9$	$p(\text{overcast} N) = 0$
$p(\text{rain} P) = 3/9$	$p(\text{rain} N) = 2/5$
temperature	
$p(\text{hot} P) = 2/9$	$p(\text{hot} N) = 2/5$
$p(\text{mild} P) = 4/9$	$p(\text{mild} N) = 2/5$
$p(\text{cool} P) = 3/9$	$p(\text{cool} N) = 1/5$
humidity	
$p(\text{high} P) = 3/9$	$p(\text{high} N) = 4/5$
$p(\text{normal} P) = 6/9$	$p(\text{normal} N) = 2/5$
windy	
$p(\text{true} P) = 3/9$	$p(\text{true} N) = 3/5$
$p(\text{false} P) = 6/9$	$p(\text{false} N) = 2/5$

Naive Bayesian Classification

- unseen sample $X = \langle \text{rain}, \text{hot}, \text{high}, \text{false} \rangle$
- $p(X|P)p(P) = p(\text{rain}|P)p(\text{hot}|P)p(\text{high}|P)p(\text{false}|P)p(P) = 3/9 * 2/9 * 3/9 * 6/9 * 9/14 = 0.010582$
- $p(X|N)p(N) = P(\text{rain}|N)P(\text{hot}|N)P(\text{high}|N)P(\text{false}|N)P(N) = 2/5 * 2/5 * 4/5 * 2/5 * 5/14 = \mathbf{0.018286}$
- Sample X is classified in class N (don't play)

Outline

- Introduction to supervised learning
- K-nn
- Naïve Bayes
- **Logistic Regression**

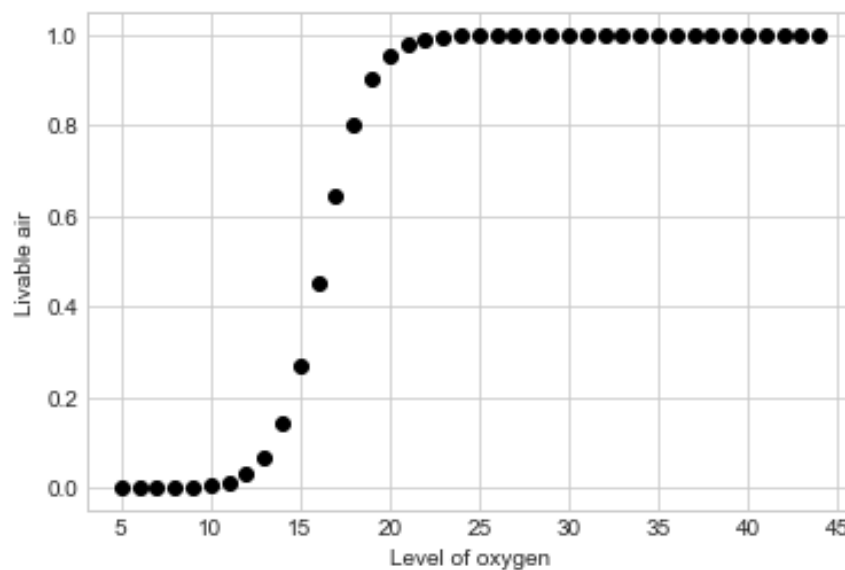
Solving a supervised learning problem

Given data, a real-world classification problem, and constraints, you need to determine:

- classifier to use
- optimization method to employ
- loss function to minimize
- features to consider from the data
- evaluation metric to use

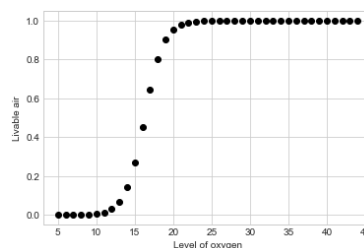
Logistic Regression

- Fits binary classification problems
- Output in $[0,1]$
- Need a function that takes data vector x and produces as output $p(x)$ in $[0,1]$



Why linear regression does not do..

- Linear regression: output y is continuous – need binary classification
- Need a $p(X)$ in $[0,1]$ – linear regression does not guarantee it.
- homoscedasticity assumption: variance of Y constant across values of X .
- significance testing assumes prediction errors $(y-f(x))$ are normally distributed.
- But y takes values 0 and 1,



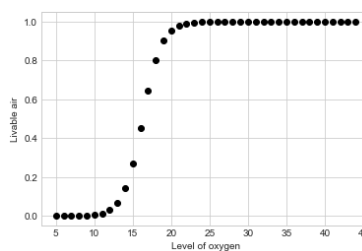
Logistic function – design

- Classification: $\mathbf{x} \in R^p, y \in \{0, 1\}$
- Assumption (odds $p(y=1)$):

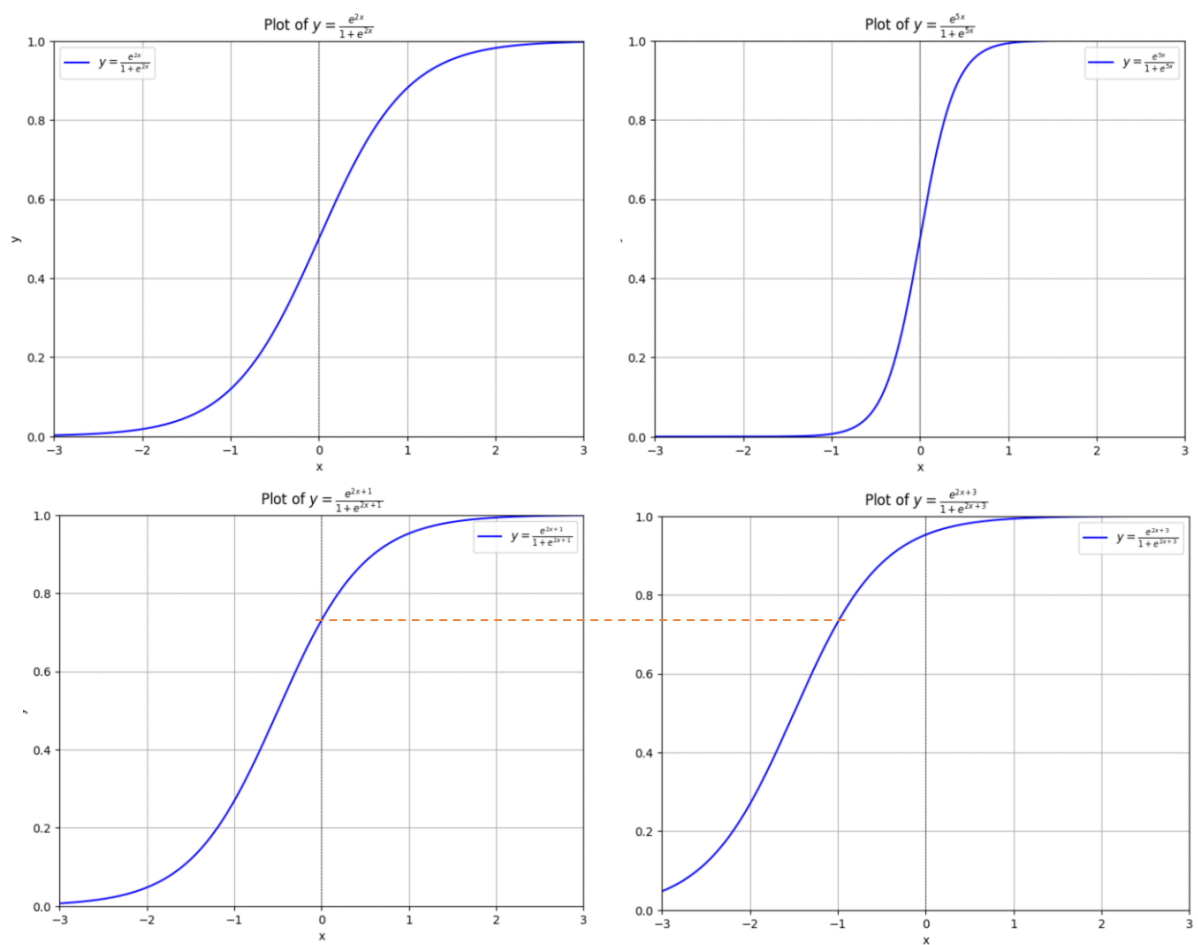
$$\log\left(\frac{p(y=1|\mathbf{x})}{1-p(y=1|\mathbf{x})}\right) = \alpha\mathbf{x} + \beta$$

- therefore

$$p(y=1|\mathbf{x}) = \frac{e^{\alpha\mathbf{x}+\beta}}{1+e^{\alpha\mathbf{x}+\beta}} = \frac{1}{1+e^{-(\alpha\mathbf{x}+\beta)}}$$



Examples of logistic function - parameters



Click prediction example

- Web site publishing ads
- Predict user i click an ad.
- Assume class c_i (1: clicked, 0: otherwise)
- x_i data for user i (history of URLs visited)
- $\text{logit}(P(\text{user } i \text{ clicks on the ad}_i)) = \text{linear function of the features } x_i$

Click Prediction: Logistic regression model

- linear model for c_i ,
- take the log of the odds ratio:

$$\log(p(c_i = 1|x_i)/(1 - p(c_i = 1|x_i))) = \alpha + \beta^T x_i$$

- Or

$$p(c_i = 1|x_i) = \frac{e^{\alpha + \beta^T x_i}}{1 + e^{\alpha + \beta^T x_i}}$$

- α : base rate, unconditional probability of $c_i = 1$ (click)
- β : dependence on the user data
 - slope of the logit function
 - Determines the relevance of features (i.e. pages visited) to the likelihood of the ad being clicked

Click Prediction: α, β parameter estimation

- Parameters: $\theta = \{\alpha, \beta\}$,
- Likelihood: $L(\theta|X_1, \dots, X_n) = p(X|\theta) = p(X_1|\theta) \dots p(X_n|\theta)$,
- X_i users are independent:

$$\theta_{MLE} = \underset{\theta}{\operatorname{argmax}} \prod_{k=1}^N p(X_i|\theta)$$

- Setting:

$$p_i = \frac{1}{1 + e^{-(\alpha + \beta^T x_i)}}$$

- Then:

$$\theta_{MLE} = \underset{\theta}{\operatorname{argmax}} \prod_{k=1}^N p_i^c (1 - p_i)^c$$

c : number of clicks

- MLE with either i. Newton-Raphson method or ii. Stochastic gradient descent